

How To Parallel Program

Shelley Knuth

shelley.knuth@colorado.edu

www.rc.colorado.edu

Questions? #RC_BasicSC

Link to survey on this topic: <http://tinyurl.com/rcpresurvey>

Slides:

https://github.com/ResearchComputing/Final_Tutorials/tree/master/Basics_Supercomputing

Outline

- Parallel Computing with OpenMP
- Parallel Computing with Matlab

Parallel Computing with Examples (OpenMP)

Outline

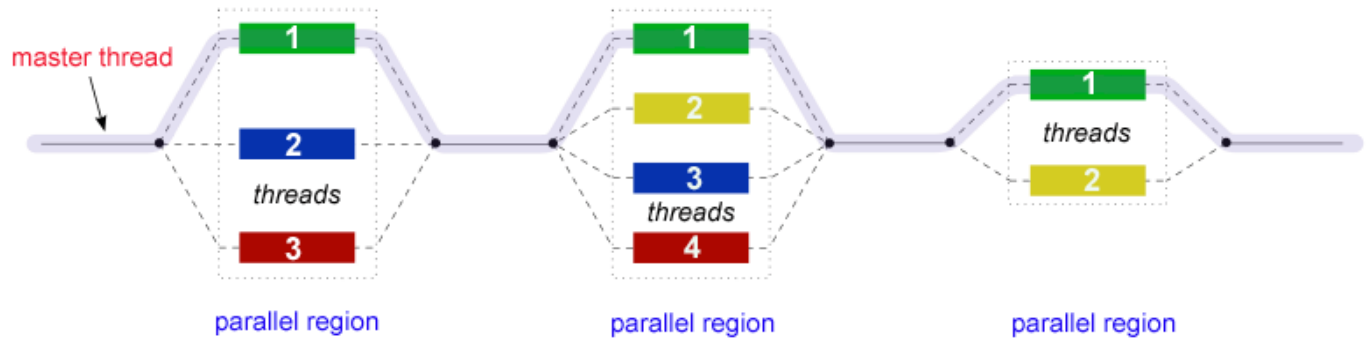
- Shared memory
- What is OpenMP?
- How is OpenMP used?
- Parallel region
- Public/Private variables
- Examples

OpenMP Directives

- Comments in source code that specify parallelism for shared memory machines
 - Enclosing parallel directives
- **FORTRAN:** directives begin with **!\$OMP**, **C\$OMP** or ***\$OMP**
- **C/C++:** directives begin with **#pragma omp**

OpenMP – Fork/Join

- OpenMP programs start with a single thread (master)
- Then Master creates a team of parallel “worker” threads (FORK)
- Statements in block are executed in parallel by every thread
- At end, all threads synchronize and join master thread



Source: <https://computing.llnl.gov/tutorials/openMP/#Introduction>

OpenMP Fortran: General Code Structure – Parallel Regions

Parallel regions are blocks of code that will be executed by multiple threads

```
1      !$OMP PARALLEL
2          code block
3          call work(...)
4      !$OMP END PARALLEL
```

Line 1	Team of threads formed at parallel region.
Lines 2-3	Each thread executes code block and subroutine calls. No branching (in or out) in a parallel region.
Line 4	All threads synchronize at end of parallel region (implied barrier).

Use the thread number to divide work among threads.

Parallel Regions

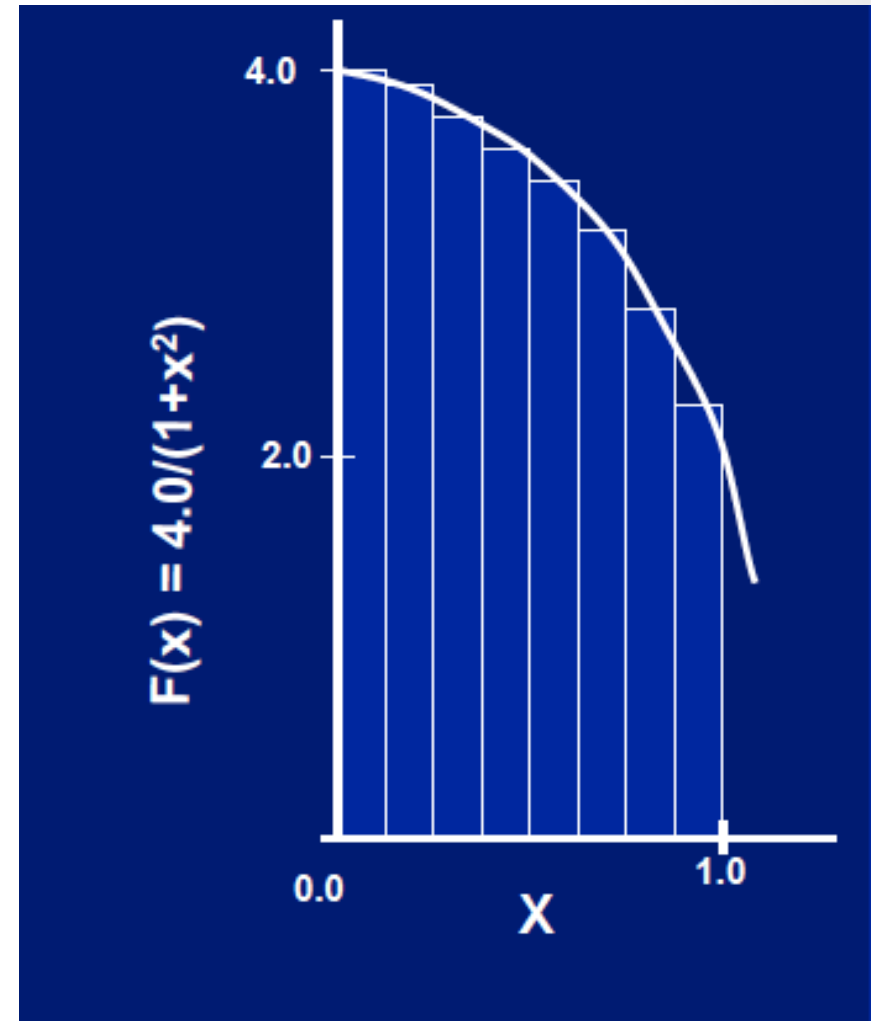
- When thread hits PARALLEL directive, creates team of threads
 - Becomes master
 - Code is duplicated and all threads execute that code
 - Runs the same code on different data
 - Split up loops and operate on different data
 - Only master thread continues after implied barrier
- Can determine number of threads by:
 - Setting the number threads to a default number or within code
 - Allowing number of threads to change from one parallel region to another

Parallel Region Example

- Finding the integral
 - Area under a curve
 - Sum of the area of all the rectangles underneath the curve (approximate)

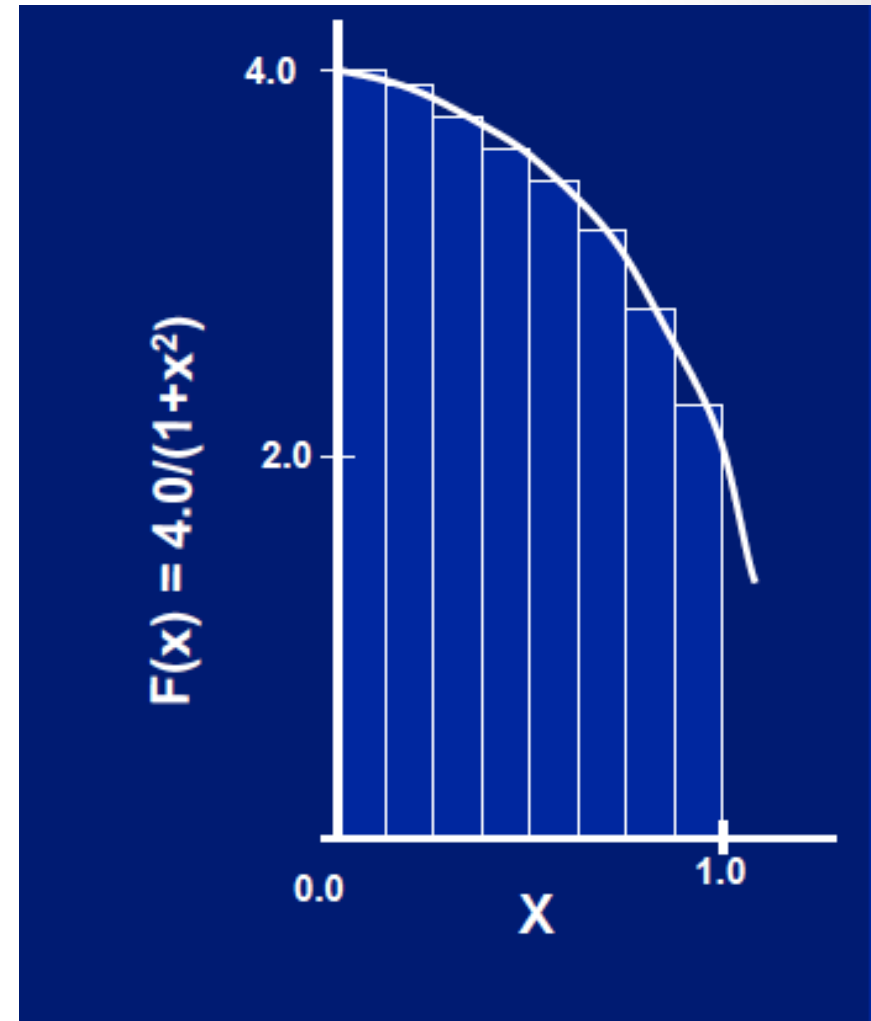
$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$

$$\sum_{i=0}^N F(x_i) \Delta x \approx \pi$$



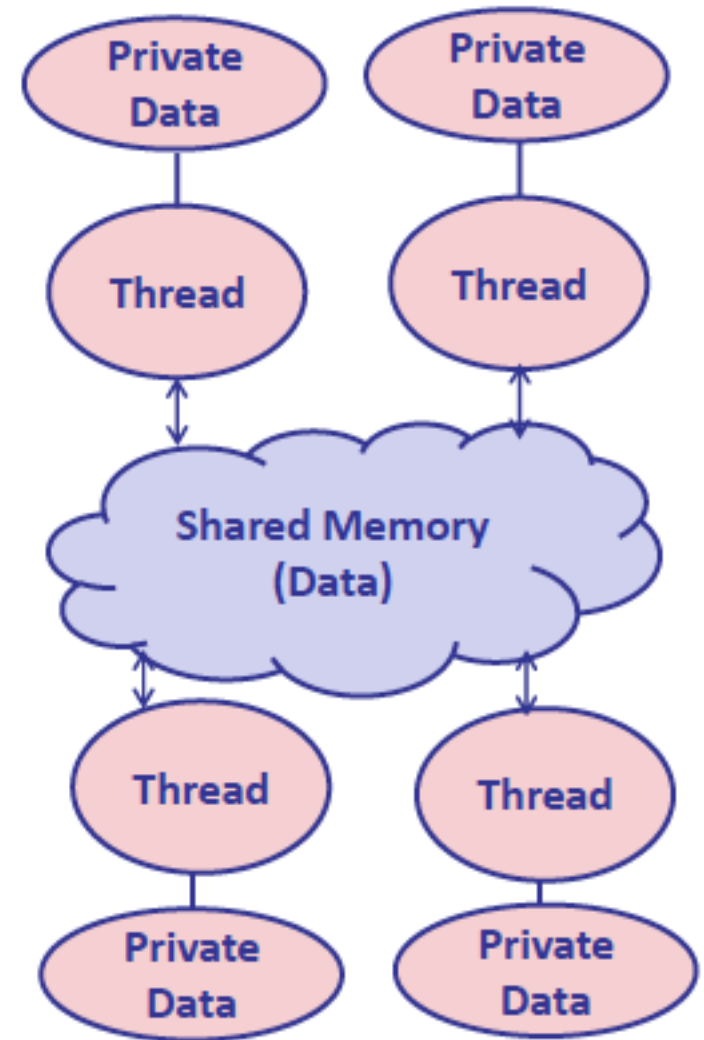
Parallel Region Example

- The same code is used to calculate the area of each of the rectangles
- Different threads will calculate different rectangles
- Which rectangles are calculated with each thread is random



Shared and Private Variables

- When specifying the **PRIVATE** clause, that variable is private to each thread
 - Each thread has own unique copy
 - Can only be accessed by the threads that own it
 - Variables declared in private subroutines are default private
 - Index variables are also default private
- When specifying **SHARED** clause, all threads can access that data
 - Global variables are shared by default



Private Variable Example

```
#pragma omp parallel for shared(a,b,c,n) private(temp,i)
    for (i=0; i< n; i++){
        temp = a[i] / b[i];
        c[i] = temp + cos(temp);
    }
```

- All threads can access a, b, c, and n
- Each loop has own private copy of index i
- Variable temp also needs to be private
- Otherwise each thread would be reading/writing to same location

Runtime Library Routines

Routine	Purpose
<u>OMP SET NUM THREADS</u>	Sets the number of threads that will be used in the next parallel region
<u>OMP GET NUM THREADS</u>	Returns the number of threads that are currently in the team executing the parallel region from which it is called
<u>OMP GET THREAD NUM</u>	Returns the thread number of the thread, within the team, making this call.
<u>OMP GET THREAD LIMIT</u>	Returns the maximum number of OpenMP threads available to a program

- In C/C++, must include the omp.h header file

Fortran	INTEGER FUNCTION OMP_GET_NUM_THREADS()
C/C++	#include <omp.h> int omp_get_num_threads(void)

OpenMP Compiling

- When compiling must use appropriate compiler flag to turn on OpenMP compilations

Compiler / Platform	Compiler	Flag
Intel Linux Opteron/Xeon	icc icpc ifort	-qopenmp
PGI Linux Opteron/Xeon	pgcc pgCC pgf77 pgf90	-mp
GNU Linux Opteron/Xeon IBM Blue Gene	gcc g++ g77 gfortran	-fopenmp
IBM Blue Gene	bgxlc_r, bgcc_r bgxlC_r, bgxlc++_r bgxlc89_r bgxlc99_r bgxlf_r bgxlf90_r bgxlf95_r bgxlf2003_r *Be sure to use a thread-safe compiler - its name ends with _r	-qsmp=omp

OMP Code Practice – Exercise 1

- Code:

```
omp_hello.f  
omp_hello.sh
```

- Instructions for running:

```
ssh tutorial-login.rc.colorado.edu -l user00XX  
ml slurm  
sbatch omp_hello.sh
```

How Do I Prepare My Code for OpenMP?

- I have code! I want it to be parallel too!
- Steps to go through
 1. Verify that code is parallelizable
 - Make sure you don't have any loop dependencies
 2. Analyze your code
 - Where does the program spend most of its time?
 - Look for loops
 - Typically easy to parallelize
 - Outside of nested loops

How Do I Prepare My Code for OpenMP?

- Steps to go through

3. Restructure code

- Put `parallel do` constructs around parallelizable loops
- List variables with appropriate `shared`, `private`, etc. clauses
- Many other things you can do that we don't cover here

4. Overhead

- How much time was spent preparing your code for parallelization?
- Is this more than the time spent running your code serially?

Example Code – Exercise 2

- Code:

```
for.c  
for.sh
```

- Instructions for running:

```
ssh tutorial-login.rc.colorado.edu -l user00XX  
ml slurm  
sbatch for.sh
```

Example Code – Exercise 2

- We need to consider whether our code really does experience a speed up
 - Array size 10,000,000
 - Drops by ~30-50%
- Let's see what happens when we change our array size to 10
 - Takes longer for parallel code to run
 - Overhead is more of a factor

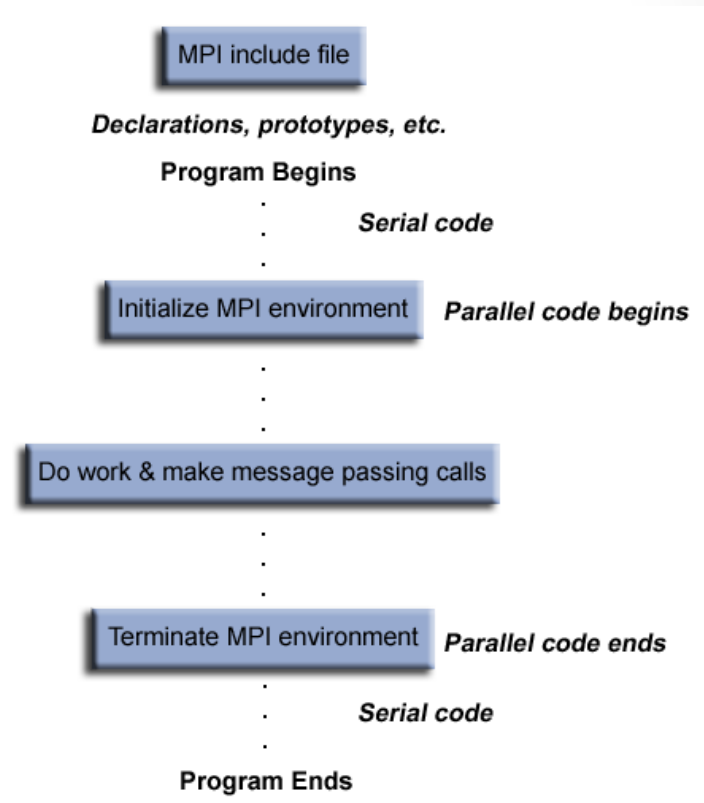
Parallel Computing with Examples (MPI)

MPI

- MPI is a library specification for message passing
- Widely used standard
- Can run on shared, distributed, or hybrid memory models
- Exchange data between processes through communication between tasks – send and receive data
- MPI can get complicated
- Programmers must explicitly implement parallelism using MPI constructs
- Portable

General MPI Code Structure

- You must have your header file at the top of any script you develop that uses MPI
- For C:
 - `#include mpi.h`
- For Fortran:
 - `use mpi`

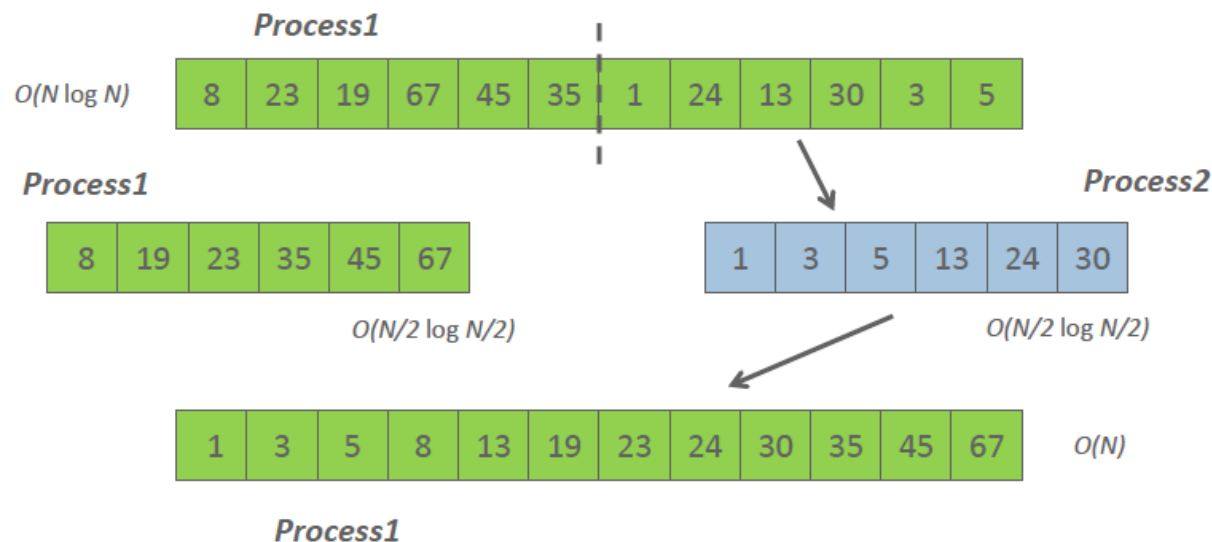


<https://computing.llnl.gov/tutorials/mpi/#What>

Message Passing

- A program that runs on a node is called a **process**
- When a program is run a process is run on each processor in the cluster
- These processes communicate with each other using message passing
- Message passing allows us to copy data from the memory of one process into another
- Message passing systems must at a minimum support system calls for sending and receiving messages

Example – Sorting Integers

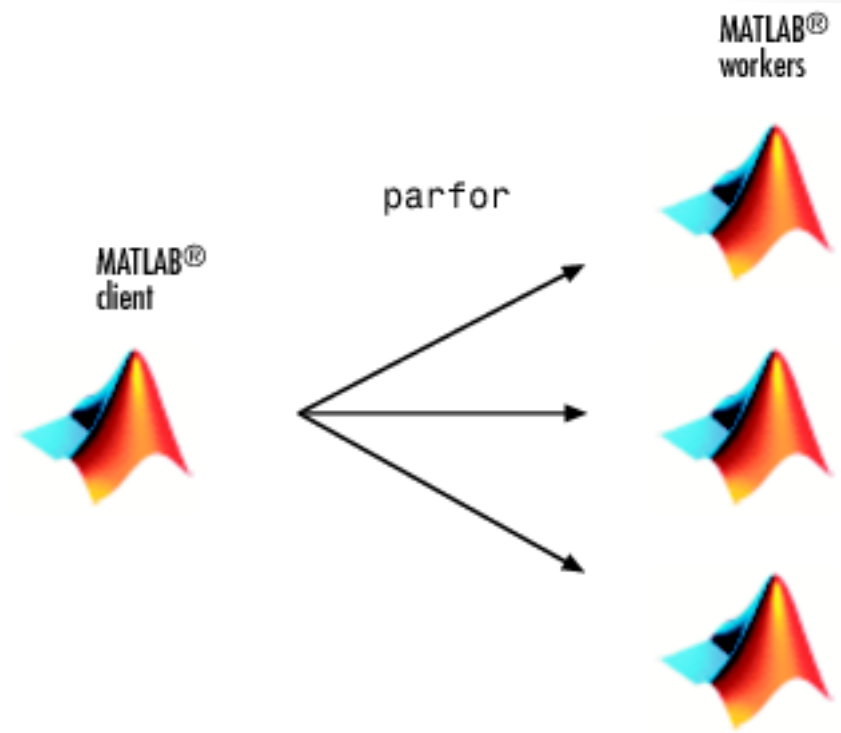


http://hlor.inf.ethz.ch/teaching/mpl_tutorials/ppopp13/2013-02-24-ppopp-mpl-basic.pdf

Parallel Computing with Examples (Matlab)

Running Matlab in Parallel

- **Workers:** copies of the original client created to assist in computation



Parallel Computing Toolbox (PCT)

- Additional toolbox as part of Matlab
- Perform parallel computations on multicore computers, GPUs, and computer clusters
- Many Matlab functions work in concert with the PCT
- Simple to utilize with just the use of certain commands

Parallel and Not Parallel

Not Parallel:

```
for i=1:10  
    x=x(i)+1;  
end
```

Parallel:

```
matlabpool open 4  
    parfor i=1:10  
        x=x(i)+1;  
    end  
matlabpool close
```

parfor

- Easy to use
- Allows parallelism in terms of loops
- When client reaches a parfor loop iterations of loop are automatically divided up among workers
- Parfor requires results be completely independent
- Cannot determine how loops are divided

Running Matlab in Parallel On Lots of Cores

- Typically see a significant speed up when using parfor vs. when not
 - If code is parallelizable
- However, this might not always be the case
- Might spend more time in overhead
 - If code isn't parallelizable
 - If code isn't that complicated

Running Matlab in Parallel

- Let's take ordinary code that is already running and convert it to run in parallel
- `matlab_parallel_serial.m`
- `matlab_parallel_tutorial.m`

Spmd Command

- Single process, multiple data
- The spmd command ensures more control
- Can parallelize much more than just loops
- Like a very simplified version of MPI
- More flexibility than parfor
- However, need to know what you're doing

Distributed Computing Toolbox

- PCT allows you to run programs in parallel across many processors
- DCT allows you to run across nodes
 - Allows you to run easily on clusters
 - Supports resource managers
 - Not an option for Janus

Other Options

- Parallel R: https://earthlab.github.io/r/R-parallel_r/
- Parallel Python:
<http://materials.jeremybejarano.com/MPIwithPython/>

References

- https://portal.tacc.utexas.edu/c/document_library/get_file?uuid=c3c38847-ca7e-41bf-aefa-fb232a777699&groupId=13601
- <https://computing.llnl.gov/tutorials/openMP/>
- <http://openmp.org/mp-documents/omp-hands-on-SC08.pdf>
- <http://heather.cs.ucdavis.edu/ParallelR.pdf>
- <https://computing.llnl.gov/tutorials/mpi/>
- http://htor.inf.ethz.ch/teaching/mpi_tutorials/ppopp13/2013-02-24-ppopp-mpi-basic.pdf
- <https://www.rc.usf.edu/tutorials/classes/tutorial/mpi/>

Now what?

- Get an account on Janus!
- Email rc-help@colorado.edu for any help!
- Join our meetup group!
<https://www.meetup.com/University-of-Colorado-Computational-Science-and-Engineering/>
- Email shelley.knuth@colorado.edu if you want to be added to a new email list about upcoming workshops!
- Fill this out! <http://tinyurl.com/curc-survey16>

Questions?

- Email rc-help@colorado.edu
- Twitter: CUBoulderRC
- Link to survey on this topic:
<http://tinyurl.com/curc-survey16>
- Slides:
https://github.com/ResearchComputing/Final_Tutorials/tree/master/Basics_Supercomputing